

Stock Return and Volatility Modeling with Earning Report Keyword Analysis

March 19, 2025

```
[1]: import os
import glob
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
import xgboost as xgb
import warnings
warnings.filterwarnings('ignore')

from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import Lasso, LinearRegression
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.feature_selection import VarianceThreshold
from datetime import datetime, timedelta
from collections import defaultdict
from math import prod
from ast import literal_eval

pd.set_option("max_colwidth", 40)
pd.options.display.max_rows = 4000
```

1 1. Classification of Keywords

In this section, I create an excel file with brand new columns for each Category in the latter section of the Data spreadsheet.

```
[2]: # Import files

txt_files = glob.glob(os.path.join("Daily_Return Data", "*.txt"))
df_data_temp = pd.read_excel(open('Data.xlsx', 'rb'))

# Filter irrelevant data as instructed
```

```

df_data = df_data_temp.copy()
df_data = df_data[df_data['documentType'] == 'EARNINGS_CALL']
df_data = df_data[df_data['country'] == 'US']

# Forward fill any missing identifiers by referencing earnings calls with the
↳ same id
# Remove event columns as instructed
# Populate nan keydriver cells with neutral score of 0
# Populate nan Method2 Category cells with neutral scores of [0,0]

df_data = df_data.sort_values(by=['companyId', 'mainIdentifier isin'])
df_data['mainIdentifier isin'] = df_data.groupby(['companyId',
↳ 'companyName'])['mainIdentifier isin'].transform(lambda x: x.ffill())

event_col = [col for col in df_data.columns if 'eventcount' in col.lower()]
df_data = df_data.drop(columns = event_col)

score_col = [col for col in df_data.columns if 'score' in col.lower()]
df_data[score_col] = df_data[score_col].fillna(0)

df_data.iloc[:, 205:] = df_data.iloc[:, 205:].fillna('[0, 0]')
sorted_m2 = sorted(list(df_data.iloc[:, 205:].columns))

```

[3]: # Group Method2 Category by prefix
Calculate the aggregate [positive,negative] score for each Category in a
↳ dynamically created column

```

def sum_tuple_group(row, cols):
    total_first = 0
    total_second = 0
    for col in cols:
        val = row[col]
        if isinstance(val, str):
            try:
                val = literal_eval(val)
            except Exception as e:
                continue
        total_first += val[0]
        total_second += val[1]
    return [total_first, total_second]

grouped_cols = defaultdict(list)
for col in sorted_m2:
    prefix = col.split(" - ")[0]
    grouped_cols[prefix].append(col)

```

```

for prefix, cols in grouped_cols.items():
    new_col_name = f"{prefix}_Total"
    df_data[new_col_name] = df_data.apply(lambda row: sum_tuple_group(row,
↪cols), axis=1)

```

```

[4]: # Export result as csv or xlsx

df_data.to_csv('question1.csv')

```

2 2. Modeling forward stock returns with keyDriver scores

Here, I merge the earning calls table with the returns table, and calculate the forward returns after each earning call.

After doing feature selection, I evaluate the performance of several traditional machine learning models.

- Reprocess the earning calls dataframe in case someone wants to start here
- Filter unnecessary documents and regions, fill in missing mainIdentifier isin
- Remove events, populate missing scores
- Add eventDate column which either will provide the same day as the earning call if the earning call is made before closing hours, or next day if the call is after hours

```

[5]: df_data2 = df_data_temp.copy()
df_data2 = df_data2[df_data2['documentType'] == 'EARNINGS_CALL']
df_data2 = df_data2[df_data2['country'] == 'US']

df_data2 = df_data2.sort_values(by=['companyId', 'mainIdentifier isin'])
df_data2['mainIdentifier isin'] = df_data2.groupby(['companyId',
↪'companyName'])['mainIdentifier isin'].transform(lambda x: x.ffill())
df_data2 = df_data2.rename(columns={'mainIdentifier isin': 'fsym_isin'})

event_col = [col for col in df_data2.columns if 'eventcount' in col.lower()]
df_data2 = df_data2.drop(columns = event_col)

score_col = [col for col in df_data2.columns if 'score' in col.lower()]
df_data2[score_col] = df_data2[score_col].fillna(0)

def roll_date(timestamp):
    dt = datetime.fromisoformat(timestamp)
    cutoff_time = dt.replace(hour=16, minute=0, second=0, microsecond=0)

    if dt >= cutoff_time:
        dt += timedelta(days=1)

    return dt.date().isoformat()

```

```
df_data2['eventDate'] = df_data2['eventTime'].apply(roll_date)
```

- Upload the returns dataset, and reorganize with pivot for each row to have its own mainIdentifier isin
- Merge the earnings and returns datasets where both returns are provided and there also a matching earning call of the same identifier

```
[6]: df_list = []
for file in txt_files:
    df = pd.read_csv(file, delimiter="\t", header=None, encoding="utf-8")
    df_list.append(df)

df_ret = pd.concat(df_list, ignore_index=True)
df_filter = df_ret[df_ret[1] != 'DATE']
df_returns = df_filter.rename(columns={0: "BENCHMARK_ID",
                                     1: "DATE",
                                     2: "SECURITY_ID",
                                     3: "Weight",
                                     4: "p_price_returns",
                                     5: "gd_class_gics_h",
                                     6: "fsym_security_perm_id",
                                     7: "p_symbol",
                                     8: "fsym_isin"})

df_returns = df_returns.sort_values(by=['fsym_isin', 'DATE'])
df_retwide = df_returns.pivot_table(index = ['fsym_isin'],
                                     columns = 'DATE',
                                     values = 'p_price_returns')

df_retwide = df_retwide.fillna(1)

df_merge = pd.merge(df_data2, df_retwide, on=['fsym_isin'], how='inner')
```

- With the eventDate provided earlier as the starting point, calculate several ranges of forward returns for each earning call

```
[7]: def oneday_ret(row):
    event_date = row['eventDate']
    future_dates = [col for col in s_rets if pd.to_datetime(col) >=
    event_date]

    if len(future_dates) > 0:
        res = future_dates[0]
        return row[res]
    else:
        return 0
```

```

def nday_ret(row, days):
    event_date = row['eventDate']
    future_dates = [col for col in s_retcoll if pd.to_datetime(col) >=
↪event_date]

    if len(future_dates) >= days:
        res = [row[col] for col in future_dates[:days]]
        cumres = (1 + pd.Series(res)).prod() - 1
        return cumres
    else:
        return None

df_merge['eventDate'] = pd.to_datetime(df_merge['eventDate'])
retcol = [col for col in df_merge.columns if col[:4].isdigit()]
retdates = pd.to_datetime(retcol)
s_retcoll = [col for _, col in sorted(zip(retdates, retcol))]

df_merge['fow_1d'] = df_merge.apply(oneday_ret, axis=1)
df_merge['fow_3d'] = df_merge.apply(lambda row: nday_ret(row, 3), axis=1)
df_merge['fow_7d'] = df_merge.apply(lambda row: nday_ret(row, 7), axis=1)

df_merge[['fow_1d', 'fow_3d', 'fow_7d']] = df_merge[['fow_1d', 'fow_3d',
↪'fow_7d']].fillna(0)
df_merge.head()

# binarize into positive return or not

df_merge['fow_1d_bin'] = df_merge['fow_1d'].apply(lambda x: 1 if x > 0 else 0)
df_merge['fow_3d_bin'] = df_merge['fow_3d'].apply(lambda x: 1 if x > 0 else 0)
df_merge['fow_7d_bin'] = df_merge['fow_7d'].apply(lambda x: 1 if x > 0 else 0)

```

Given that we are looking at a dataset with many features and also happens to be very sparse, there are a few steps we should take. - We start by removing highly correlated features as a best practice. - Next, as a blanket approach to filtering sparsely populated features, we can filter by each feature's variance below a threshold, removing features that have little change and therefore little contribution. - Scale the independent variables.

There are a few models to try with sparse data: - Linear Regression as a baseline check - Lasso Regularization to penalize certain features - RandomForest and XGBoost, as decision trees are well-equipped for sparse data

We'll evaluate the performance of each model by looking at their R^2 value, and attempt to identify reoccurring features with high importance to see if there are any that can be used to predict positive forward returns.

```
[8]: # remove correlated features

drive_col = [col for col in df_merge.columns if 'keydriver' in col.lower()]

df_corr = df_merge[drive_col]
mat = df_corr.corr()

to_remove = set()
for i in range(len(mat.columns)):
    for j in range(i):
        if abs(mat.iloc[i, j]) > 0.9:
            colname = mat.columns[i]
            to_remove.add(colname)

print(to_remove)

df_tpp = df_merge[drive_col].drop(columns=to_remove)
```

```
{'Qna Deception keyDriver positiveScore', 'Total Exec Change keyDriver
positiveScore', 'Answer Exec Change keyDriver positiveScore', 'Answer Deception
keyDriver negativeScore', 'Answer Guidance keyDriver positiveScore', 'Total
Capital Raise Returns keyDriver positiveScore', 'Answer Headwinds Tailwinds
keyDriver negativeScore', 'Total Exec Change keyDriver negativeScore', 'Total
Wage keyDriver score'}
```

- Due to there being many features, some features may be very sparse and be mostly 0s. In that event, we filter and remove the features with low variance.

```
[9]: selector = VarianceThreshold(threshold=0.01) # Remove features with very low
↪variance
df_reduced = df_tpp.iloc[:, selector.fit(df_tpp).get_support()]
```

2.1 Linear Regression

```
[10]: X = df_reduced
feature_names = df_reduced.columns

#y1d = df_merge['fow_1d']
#y3d = df_merge['fow_3d']
#y7d = df_merge['fow_7d']

y1d = df_merge['fow_1d_bin']
y3d = df_merge['fow_3d_bin']
y7d = df_merge['fow_7d_bin']
```

```
[11]: forwards = [(1, y1d), (3, y3d), (5, y7d)]

for num, table in forwards:
```

```

X_train, X_test, y_train, y_test = train_test_split(X, table, test_size=0.
↪2, random_state=42)

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

feature_importance = np.abs(model.coef_)
df_feat = pd.DataFrame({'Feature': feature_names, 'Importance': ↪
↪feature_importance})
df_feat = df_feat.sort_values(by="Importance", ascending=False)

print(f"{num}-Day Forward Return Linear Regression")
print("R^2:", r2)
print("MSE:", mse)
print(df_feat.head(10))
print(" ")

```

1-Day Forward Return Linear Regression

R^2: -0.21224036354801634

MSE: 0.30305084025085477

	Feature	Importance
136	Presentation Irregularities keyDrive...	1.289802
130	Total Irregularities keyDriver score	1.115469
139	Answer Merger Acquisition keyDriver ...	0.694026
137	Presentation Irregularities keyDrive...	0.621898
131	Total Irregularities keyDriver negat...	0.536375
145	Qna Exec Change keyDriver score	0.490342
146	Qna Exec Change keyDriver negativeScore	0.483078
138	Presentation Irregularities keyDrive...	0.447071
133	Answer Irregularities keyDriver score	0.432729
124	Answer Capital Raise Returns keyDriv...	0.385358

3-Day Forward Return Linear Regression

R^2: -0.5255556743718135

MSE: 0.3598636870530623

	Feature	Importance
145	Qna Exec Change keyDriver score	1.323129
139	Answer Merger Acquisition keyDriver ...	1.227862
133	Answer Irregularities keyDriver score	1.031828
146	Qna Exec Change keyDriver negativeScore	0.835097
135	Answer Irregularities keyDriver posi...	0.791380

```

147 Qna Exec Change keyDriver positiveScore    0.567497
142 Question Irregularities keyDriver score    0.565293
127     Qna Irregularities keyDriver score    0.524743
140 Answer Merger Acquisition keyDriver ...    0.505078
143 Question Irregularities keyDriver ne...    0.446560

```

5-Day Forward Return Linear Regression

R^2 : -0.3829553723353947

MSE: 0.28637615597580396

	Feature	Importance
145	Qna Exec Change keyDriver score	1.406079
133	Answer Irregularities keyDriver score	1.163594
146	Qna Exec Change keyDriver negativeScore	0.991046
135	Answer Irregularities keyDriver posi...	0.644708
134	Answer Irregularities keyDriver nega...	0.614579
139	Answer Merger Acquisition keyDriver ...	0.596427
125	Answer Capital Raise Returns keyDriv...	0.466832
147	Qna Exec Change keyDriver positiveScore	0.450984
141	Answer Merger Acquisition keyDriver ...	0.403433
137	Presentation Irregularities keyDrive...	0.400318

3 Lasso Regularization

```

[12]: for num, table in forwards:
    X_train, X_test, y_train, y_test = train_test_split(X, table, test_size=0.
↪2, random_state=42)

    param_grid = {'alpha': np.logspace(-4, 1, 50)}

    lasso = Lasso()

    grid_search = GridSearchCV(lasso, param_grid, cv=5,
↪scoring='neg_mean_squared_error')
    grid_search.fit(X_train, y_train)

    best_alpha = grid_search.best_params_['alpha']

    lasso_best = Lasso(alpha=best_alpha)
    lasso_best.fit(X_train, y_train)
    y_pred = lasso_best.predict(X_test)

    r2 = r2_score(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)

    feature_importance = np.abs(lasso_best.coef_)

```



```

df_las_imp = pd.DataFrame({'Feature': feature_names, 'Importance': ↵
↵feature_importance})
df_las_imp = df_las_imp.sort_values(by="Importance", ascending=False)

print(f"{num}-Day Forward Return LassoCV")
print("R^2:", r2)
print("MSE:", mse)
print(df_las_imp.head(10))
print(" ")

```

1-Day Forward Return LassoCV

R²: 0.02413155549624013

MSE: 0.24395966424974183

	Feature	Importance
10	Qna Deception keyDriver negativeScore	0.003581
29	Answer Deception keyDriver score	0.000036
31	Question Capital Raise Returns keyDr...	0.000027
97	Presentation Headwinds Tailwinds key...	0.000000
98	Presentation Headwinds Tailwinds key...	0.000000
99	Answer Market Position keyDriver score	0.000000
100	Answer Market Position keyDriver neg...	0.000000
101	Answer Market Position keyDriver pos...	0.000000
102	Question Margin keyDriver score	0.000000
103	Question Margin keyDriver negativeScore	0.000000

3-Day Forward Return LassoCV

R²: -0.015915358271915148

MSE: 0.23964451294909678

	Feature	Importance
36	Presentation Wage keyDriver positive...	0.006740
66	Total Headwinds Tailwinds keyDriver ...	0.005899
65	Total Headwinds Tailwinds keyDriver ...	0.004715
10	Qna Deception keyDriver negativeScore	0.003642
18	Total Merger Acquisition keyDriver p...	0.003199
14	Total Deception keyDriver negativeScore	0.001934
29	Answer Deception keyDriver score	0.000027
31	Question Capital Raise Returns keyDr...	0.000025
0	Qna Wage keyDriver score	0.000000
100	Answer Market Position keyDriver neg...	0.000000

5-Day Forward Return LassoCV

R²: -0.005223245264588838

MSE: 0.20815709214843767

	Feature	Importance
29	Answer Deception keyDriver score	0.000044
31	Question Capital Raise Returns keyDr...	0.000024
104	Question Margin keyDriver positiveScore	0.000000

```

97 Presentation Headwinds Tailwinds key... 0.000000
98 Presentation Headwinds Tailwinds key... 0.000000
99 Answer Market Position keyDriver score 0.000000
100 Answer Market Position keyDriver neg... 0.000000
101 Answer Market Position keyDriver pos... 0.000000
102 Question Margin keyDriver score 0.000000
103 Question Margin keyDriver negativeScore 0.000000

```

3.1 Random Forest

```

[13]: for num, table in forwards:
        X_train, X_test, y_train, y_test = train_test_split(X, table, test_size=0.
        ↪2, random_state=42)

        rf = RandomForestRegressor(n_estimators=100, random_state=42)
        rf.fit(X_train, y_train)

        y_pred = rf.predict(X_test)

        r2 = r2_score(y_test, y_pred)
        mse = mean_squared_error(y_test, y_pred)

        df_grad_imp = pd.DataFrame({"Feature": X.columns, "Importance": rf.
        ↪feature_importances_})
        df_grad_imp = df_grad_imp.sort_values(by="Importance", ascending=False)

        print(f"{num}-Day Forward Return Random Forest")
        print("R^2:", r2)
        print("MSE:", mse)
        print(df_grad_imp.head(10))
        print(" ")

```

1-Day Forward Return Random Forest

R²: -0.05129175824175847

MSE: 0.26281491712707183

	Feature	Importance
9	Qna Deception keyDriver score	0.030538
19	Total Guidance keyDriver score	0.025018
22	Total Capital Raise Returns keyDrive...	0.021717
16	Total Merger Acquisition keyDriver s...	0.020661
29	Answer Deception keyDriver score	0.020611
27	Answer Guidance keyDriver score	0.019777
37	Presentation Deception keyDriver score	0.019405
52	Qna Market Position keyDriver score	0.019033
87	Presentation Guidance keyDriver score	0.018814
96	Presentation Headwinds Tailwinds key...	0.018164

3-Day Forward Return Random Forest

R²: -0.06820681935817796

MSE: 0.2519795580110497

	Feature	Importance
9	Qna Deception keyDriver score	0.026530
10	Qna Deception keyDriver negativeScore	0.023715
61	Total Market Position keyDriver score	0.023457
64	Total Headwinds Tailwinds keyDriver ...	0.022962
96	Presentation Headwinds Tailwinds key...	0.021703
22	Total Capital Raise Returns keyDrive...	0.019802
19	Total Guidance keyDriver score	0.019338
34	Presentation Wage keyDriver score	0.018937
46	Qna CapEx keyDriver score	0.018712
27	Answer Guidance keyDriver score	0.018705

5-Day Forward Return Random Forest

R²: -0.051718322523585325

MSE: 0.21778508287292817

	Feature	Importance
27	Answer Guidance keyDriver score	0.030829
96	Presentation Headwinds Tailwinds key...	0.024310
64	Total Headwinds Tailwinds keyDriver ...	0.021355
29	Answer Deception keyDriver score	0.020693
3	Qna Guidance keyDriver score	0.019774
9	Qna Deception keyDriver score	0.019415
31	Question Capital Raise Returns keyDr...	0.019271
16	Total Merger Acquisition keyDriver s...	0.018800
19	Total Guidance keyDriver score	0.018421
22	Total Capital Raise Returns keyDrive...	0.018031

3.2 XGBoost Model

```
[14]: for num, table in forwards:
```

```
    X_train, X_test, y_train, y_test = train_test_split(X, table, test_size=0.
↪2, random_state=42)
```

```
    param_grid = {
        "n_estimators": [100, 200],
        "learning_rate": [0.01, 0.1],
        "max_depth": [3, 5, 7],
        "subsample": [0.8, 1.0]}
```

```
    xgb_model = xgb.XGBRegressor(objective="reg:squarederror", random_state=42)
```

```

grid_search = GridSearchCV(xgb_model, param_grid, cv=3, scoring="r2",
↳ verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)

best_xgb = grid_search.best_estimator_
best_params = grid_search.best_params_
y_pred = best_xgb.predict(X_test)

r2_xgb = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

df_xgb_imp = pd.DataFrame({"Feature": X.columns, "Importance": best_xgb.
↳ feature_importances_})
df_xgb_imp = df_xgb_imp.sort_values(by="Importance", ascending=False)

print(f"{num}-Day Forward Return LassoCV")
print("R^2:", r2_xgb)
print("MSE:", mse)
print(df_xgb_imp.head(10))
print(" ")

```

Fitting 3 folds for each of 24 candidates, totalling 72 fits

1-Day Forward Return LassoCV

R^2: -0.008499583855449577

MSE: 0.2521172000786341

	Feature	Importance
107	Question Headwinds Tailwinds keyDriv...	0.014660
112	Answer CapEx keyDriver negativeScore	0.012348
39	Presentation Deception keyDriver pos...	0.011403
117	Total Exec Change keyDriver score	0.010794
37	Presentation Deception keyDriver score	0.010672
54	Qna Market Position keyDriver positi...	0.010520
47	Qna CapEx keyDriver negativeScore	0.010519
96	Presentation Headwinds Tailwinds key...	0.010306
83	Presentation CapEx keyDriver positiv...	0.010299
52	Qna Market Position keyDriver score	0.010230

Fitting 3 folds for each of 24 candidates, totalling 72 fits

3-Day Forward Return LassoCV

R^2: -0.052880535815356566

MSE: 0.24836423737923372

	Feature	Importance
106	Question Headwinds Tailwinds keyDriv...	0.010688
30	Answer Deception keyDriver positiveS...	0.009308
39	Presentation Deception keyDriver pos...	0.009302
120	Question Merger Acquisition keyDrive...	0.009234
119	Question Merger Acquisition keyDrive...	0.009126

95	Presentation Capital Raise Returns k...	0.009043
62	Total Market Position keyDriver nega...	0.009026
122	Presentation Exec Change keyDriver n...	0.008993
14	Total Deception keyDriver negativeScore	0.008854
96	Presentation Headwinds Tailwinds key...	0.008813

Fitting 3 folds for each of 24 candidates, totalling 72 fits

5-Day Forward Return LassoCV

R²: -0.011655558543950395

MSE: 0.209489066547485

	Feature	Importance
56	Total CapEx keyDriver negativeScore	0.015138
57	Total CapEx keyDriver positiveScore	0.014905
43	Qna Margin keyDriver score	0.013694
45	Qna Margin keyDriver positiveScore	0.013383
99	Answer Market Position keyDriver score	0.013108
17	Total Merger Acquisition keyDriver n...	0.013046
84	Presentation Margin keyDriver score	0.013036
46	Qna CapEx keyDriver score	0.012995
65	Total Headwinds Tailwinds keyDriver ...	0.012912
109	Question Guidance keyDriver negative...	0.012886

3.3 Part 2 Reflection

- Traditional machine learning models ineffective at capturing the relationship between key-driver feature scores and forward returns, demonstrated by the low R² values across all models.
- At best, the LassoCV regularization model returns positive R² values for the single-day forward return range.
- The overall keyDriver feature importance is very low for most models, with the only exceptions being seen in the Linear Regression model, with some of the most prevalent being ‘Presentation Irregularities’, and ‘QNA Exec Change’
- If I had more time or started from scratch, I would experiment with approaches using deep learning models, which may be able to find the relationship between features and returns more effectively.

4 3. Modeling forward stock returns with constructedcategories

- For this section, I organize, preprocess, and evaluate the data with a variety of models.

```
[15]: # Filter irrelevant data as instructed

df_data = df_data_temp.copy()
df_data = df_data[df_data['documentType'] == 'EARNINGS_CALL']
df_data = df_data[df_data['country'] == 'US']
```

```

# Forward fill any missing identifiers by referencing earnings calls with the
↳ same id
# Remove event columns as instructed
# Populate nan keydriver cells with neutral score of 0
# Populate nan Method2 Category cells with neutral scores of [0,0]

df_data = df_data.sort_values(by=['companyId', 'mainIdentifier isin'])
df_data['mainIdentifier isin'] = df_data.groupby(['companyId',
↳ 'companyName'])['mainIdentifier isin'].transform(lambda x: x.ffill())
df_data = df_data.rename(columns={'mainIdentifier isin': 'fsym_isin'})

event_col = [col for col in df_data.columns if 'eventcount' in col.lower()]
df_data = df_data.drop(columns = event_col)

score_col = [col for col in df_data.columns if 'score' in col.lower()]
df_data = df_data.drop(columns = score_col)

df_data.iloc[:, 28:]
df_data.iloc[:, 28:] = df_data.iloc[:, 28:].fillna('[0, 0]')

# Group Method2 Category by prefix
# Calculate the aggregate [positive,negative] score for each Category in a
↳ dynamically created column

for prefix, cols in grouped_cols.items():
    new_col_name = f"{prefix}_Total"
    df_data[new_col_name] = df_data.apply(lambda row: sum_tuple_group(row,
↳ cols), axis=1)

cats = list(df_data.iloc[:,28:].columns)

```

- Since [0,0] is a very inconvenient string to work with, I create two new columns dynamically
- One column will be the positive score while the other depicts the negative score.
- Then merge with the returns dataset to filter out earnings without returns or unassociated returns

```

[16]: def safe_parse_list(value):
    try:
        if isinstance(value, str) and value.startswith("[") and value.
↳ endswith("]"):
            return literal_eval(value)
        else:
            return [0, 0]
    except (ValueError, SyntaxError):
        return [0, 0]

```

```

for col in cats:
    df_data[col] = df_data[col].apply(safe_parse_list)

    df_data[[f"{col} pos", f"{col} neg"]] = pd.DataFrame(df_data[col].tolist(),
↳index=df_data.index)
    df_data.drop(columns=[col], inplace=True)

df_merge3 = pd.merge(df_data, df_retwide, on=['fsym_isin'], how='inner')

df_inputs3 = df_merge3.iloc[:,28:770]

```

- Use correlation matrix to remove highly correlated variables, although the quantity of these is very low.
- Since these categories are very sparse, the variance threshold is effective at removing the majority of categories that are mostly 0s.

```

[17]: mat = df_inputs3.corr()

to_remove = set()
for i in range(len(mat.columns)):
    for j in range(i):
        if abs(mat.iloc[i, j]) > 0.9:
            colname = mat.columns[i]
            to_remove.add(colname)

print(to_remove)

df_tepp = df_inputs3.drop(columns=to_remove)

selector = VarianceThreshold(threshold=0.01) # Remove features with very low
↳variance
df_rduced = df_tepp.iloc[:, selector.fit(df_tepp).get_support()]

```

```
{'Employment - Hypothetical neg', 'Spin Off Split Off neg'}
```

4.1 Linear Regression

```

[18]: X = df_rduced
feature_names = df_rduced.columns

#y1d = df_merge['fow_1d']
#y3d = df_merge['fow_3d']
#y7d = df_merge['fow_7d']

y1d = df_merge['fow_1d_bin']
y3d = df_merge['fow_3d_bin']
y7d = df_merge['fow_7d_bin']

```

```

for num, table in forwards:
    X_train, X_test, y_train, y_test = train_test_split(X, table, test_size=0.
↪2, random_state=42)

    model = LinearRegression()
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

    r2 = r2_score(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)

    feature_importance = np.abs(model.coef_)
    df_feat = pd.DataFrame({'Feature': feature_names, 'Importance':
↪feature_importance})
    df_feat = df_feat.sort_values(by="Importance", ascending=False)

    print(f"{num}-Day Forward Return Linear Regression")
    print("R^2:", r2)
    print("MSE:", mse)
    print(df_feat.head(10))
    print(" ")

```

1-Day Forward Return Linear Regression

R^2: -0.26443322158008

MSE: 0.31609865647388213

	Feature	Importance
225	Product Trial Results pos	0.485794
236	Settlement - Question neg	0.439078
212	Employment neg	0.411603
240	Dividend - Question pos	0.325062
216	Customer Spending - Hypothetical neg	0.308125
73	Tailwinds neg	0.280274
196	Writedowns Impairments neg	0.239461
217	Merger Acquisition Announcement pos	0.236729
172	Financial Results - Question neg	0.204233
181	Volume - Qualifier neg	0.202834

3-Day Forward Return Linear Regression

R^2: -0.56758924101893

MSE: 0.36977899498166394

	Feature	Importance
225	Product Trial Results pos	0.346474
216	Customer Spending - Hypothetical neg	0.339190
227	Debt Financing - Forecast pos	0.331088

221	Product Approval	pos	0.317388
150	Competition	pos	0.307625
208	Lawsuit	neg	0.306906
67	Strategic Alliance	neg	0.301037
154	Commodity Price - Hypothetical	pos	0.272337
243	Facilities - Qualifier	neg	0.262748
212	Employment	neg	0.236738

5-Day Forward Return Linear Regression

R²: -0.9583001858064479

MSE: 0.4055159628982917

	Feature	Importance
244	Catastrophe Loss neg	0.556881
227	Debt Financing - Forecast pos	0.539186
215	Customer Traffic - Question neg	0.387171
212	Employment neg	0.379476
243	Facilities - Qualifier neg	0.312463
223	Merger Acquisition - Hypothetical neg	0.309424
235	Settlement - Question pos	0.303110
154	Commodity Price - Hypothetical pos	0.289330
137	Margin Commentary - Question pos	0.268737
221	Product Approval pos	0.266522

4.2 LassoCV

```
[19]: for num, table in forwards:
    X_train, X_test, y_train, y_test = train_test_split(X, table, test_size=0.
    ↪2, random_state=42)

    param_grid = {'alpha': np.logspace(-4, 1, 50)}

    lasso = Lasso()

    grid_search = GridSearchCV(lasso, param_grid, cv=5,
    ↪scoring='neg_mean_squared_error')
    grid_search.fit(X_train, y_train)

    best_alpha = grid_search.best_params_['alpha']

    lasso_best = Lasso(alpha=best_alpha)
    lasso_best.fit(X_train, y_train)
    y_pred = lasso_best.predict(X_test)

    r2 = r2_score(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
```

```

feature_importance = np.abs(lasso_best.coef_)

df_las_imp = pd.DataFrame({'Feature': feature_names, 'Importance':  
↪feature_importance})
df_las_imp = df_las_imp.sort_values(by="Importance", ascending=False)

print(f"{num}-Day Forward Return LassoCV")
print("R^2:", r2)
print("MSE:", mse)
print(df_las_imp.head(10))
print(" ")

```

1-Day Forward Return LassoCV

R²: 0.0069752908721810725

MSE: 0.2482485994858776

	Feature	Importance
33	Business Commentary - Forecast neg	0.009162
2	Financial Commentary pos	0.009026
3	Financial Commentary neg	0.008282
15	AmenityQuestionTopic neg	0.006433
7	Financial Results neg	0.006395
42	Category Commentary pos	0.004207
9	Euphemism neg	0.002101
26	Margin Commentary pos	0.002028
17	Business Commentary neg	0.001393
16	Business Commentary pos	0.000853

3-Day Forward Return LassoCV

R²: -0.015382980445472594

MSE: 0.23951893021832701

	Feature	Importance
115	Weather neg	0.031161
58	Capacity Production pos	0.028664
80	Price - Question pos	0.020368
5	Facilities neg	0.017516
17	Business Commentary neg	0.015964
8	Euphemism pos	0.014641
64	Price - Forecast pos	0.013334
68	Financial Commentary - Qualifier pos	0.012982
72	Tailwinds pos	0.012699
19	Financial Commentary - Forecast neg	0.012107

5-Day Forward Return LassoCV

R²: -0.0040628685141519405

MSE: 0.20791680656878614

	Feature	Importance
0	Price pos	0.0
154	Commodity Price - Hypothetical pos	0.0

156	Price - Hypothetical neg	0.0
157	Cost - Hypothetical pos	0.0
158	Cost - Hypothetical neg	0.0
159	Consumer Trend pos	0.0
160	Consumer Trend neg	0.0
161	Tax - Forecast pos	0.0
162	Tax - Forecast neg	0.0
163	Tax - Question neg	0.0

4.3 Part 3 Reflection

- Just as in Part 2, traditional machine learning models experience difficulty in finding the relationship between features and positive returns, with the LassoCV regularization model offering the best R^2 value for a single-day future.
- Regardless, the linear model identified some of the most important features to be ‘Product Trials - Positive’, ‘Customer Spending - Negative’, and ‘Catastrophic Loss - Negative’.
- Contextually, these features make sense. A successful product trial, being the cornerstone of a company is an essential part of success. Meanwhile, the latter two cases highlight the ramifications of two negative occurrences to a company’s earnings.

[]: